



# Neural Network Basics

*FA690 Machine Learning in Finance*

**Dr. Zonghao Yang**

2025 Spring

# Learning Objectives


- Describe the architecture of a neural network, including input, hidden, and output layers, and explain the role of activation functions such as ReLU, sigmoid, and softmax in enabling nonlinear decision boundaries.
- Understand the principles of gradient descent, including backpropagation and optimization techniques such as momentum and adaptive learning rates.
- Identify key hyperparameters of a neural network, such as learning rate, batch size, and regularization techniques (dropout, L2 penalty), and develop intuition for tuning them based on learning curves and validation performance.



# Fundamentals of Neural Networks


# Introduction to the Perceptron

The base architecture




# Introduction to the Perceptron

The base architecture




# Introduction to the Perceptron

The base architecture




# Introduction to the Perceptron

The base architecture



# Introduction to the Perceptron

## The base architecture



Linear combination of inputs

$$\hat{y} = w_0 + \sum_i w_i x_i$$



Bias (intercept)

So far, this resembles a linear regression





# Introduction to the Perceptron

But how could we separate the red and green points?



# Introduction to the Perceptron

But how could we separate the red and green points?

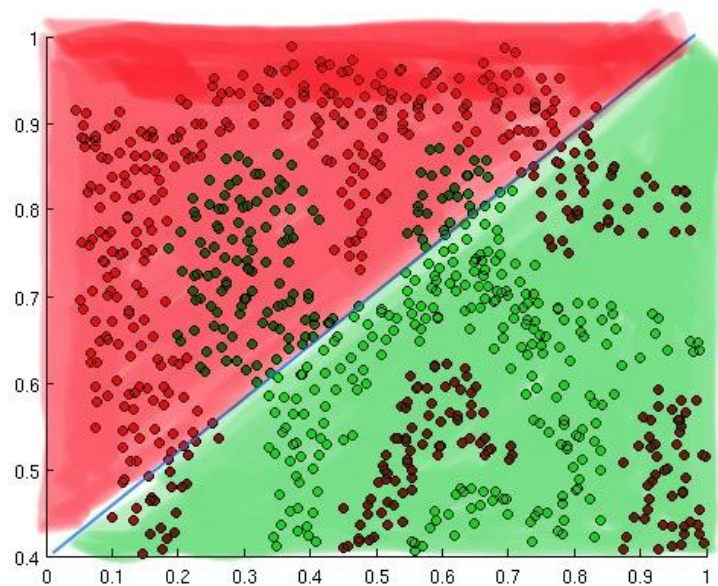


Linear functions result in linear decision boundaries, regardless of network size


# Introduction to the Perceptron

Why (nonlinear) activation functions?

$$z = w_0 + \sum_i w_i x_i \longrightarrow \hat{y} = g(z), \text{ where } g \text{ is a non-linear activation function.}$$



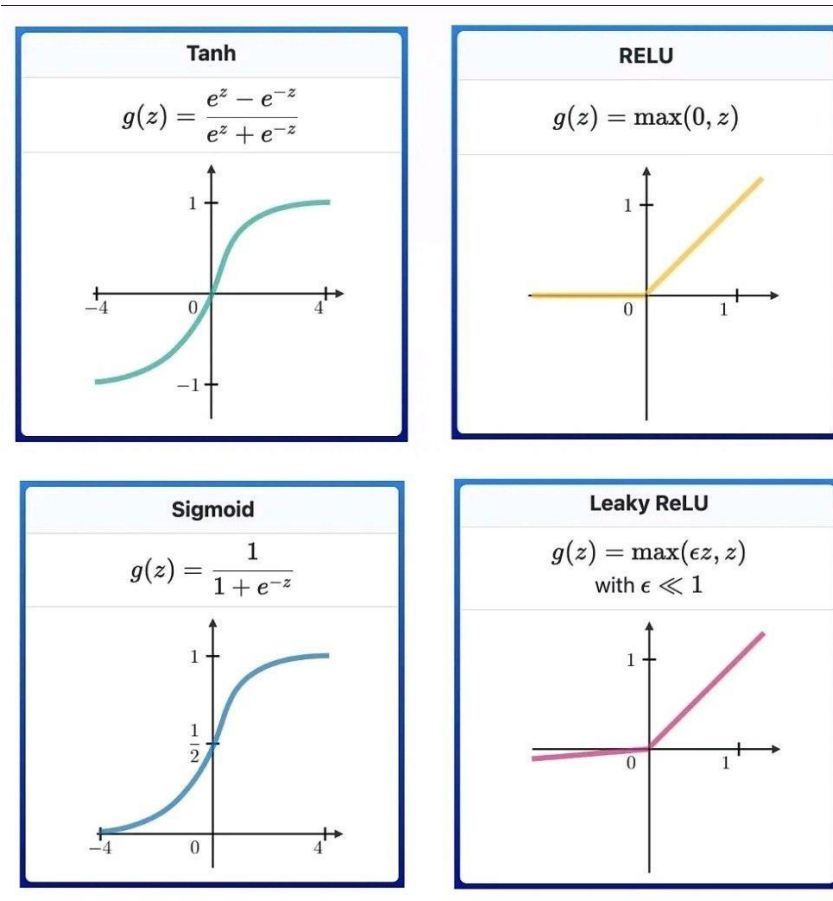
Linear functions result in linear decision boundaries, regardless of network size



Nonlinear activation functions enable the approximation of arbitrarily complex functions (nonlinear decision boundaries)


# Introduction to the Perceptron

## Classic activation functions





# Introduction to the Perceptron

The full perceptron: forward propagation




# Introduction to the Perceptron

## Artificial Neuron vs Brain Neuron



# Neural Networks

## Multi-output perceptron




All inputs are densely connected to the outputs. We call this neural network structure a “dense layer.”



# Neural Networks

## Multi-output perceptron




All inputs are densely connected to the outputs. We call this neural network structure a “dense layer.”




# Neural Networks

Single hidden layer network



# Neural Networks

## Single hidden layer network




$$z_2 = w_{0,2}^1 + \sum_i w_{i,2}^1 x_i$$

$$z_2 = w_{0,2}^1 + w_{1,2}^1 x_1 + w_{2,2}^1 x_2 + \dots + w_{n,2}^1 x_n$$


# Neural Networks

Single hidden layer network



# Neural Networks

Single hidden layer network



# Activation Functions for the Output Layer

- Regression task: The final layer is just a single neuron with no activation

$$y = g(z) = z$$

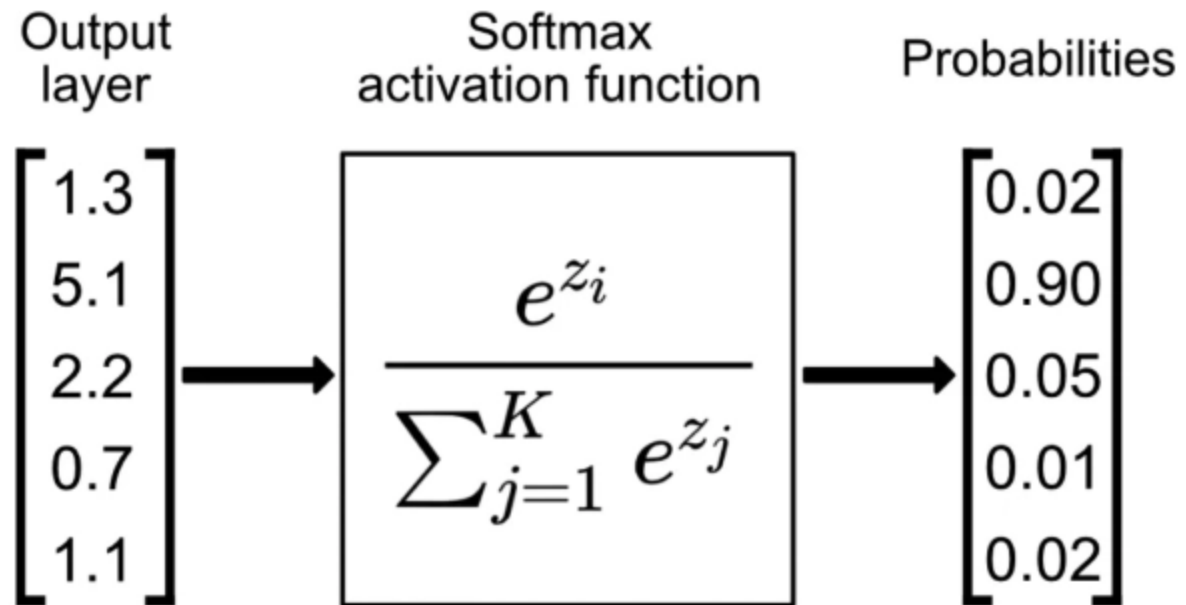
- Classification task with 2 classes: The final layer is a single neuron with the sigmoid activation function

$$y = g(z) = \text{sigmoid}(z) = \frac{1}{1 + e^{-z}}$$


- Classification task with  $K$  classes: The final layer is composed of  $K$  neurons and a softmax function that combines their outputs through normalization

$$y_i = g(z_i) = \text{softmax}(z_i) = \frac{1 + e^{-z_i}}{\sum_{j=1}^K 1 + e^{-z_j}}, i = \{1, \dots, K\}$$

# Softmax



# Deep Neural Network



$$z_{k,j} = w_{0,j}^k + \sum_{i=1}^{d_{k-1}} w_{i,j}^k g(z_{k-1,i})$$

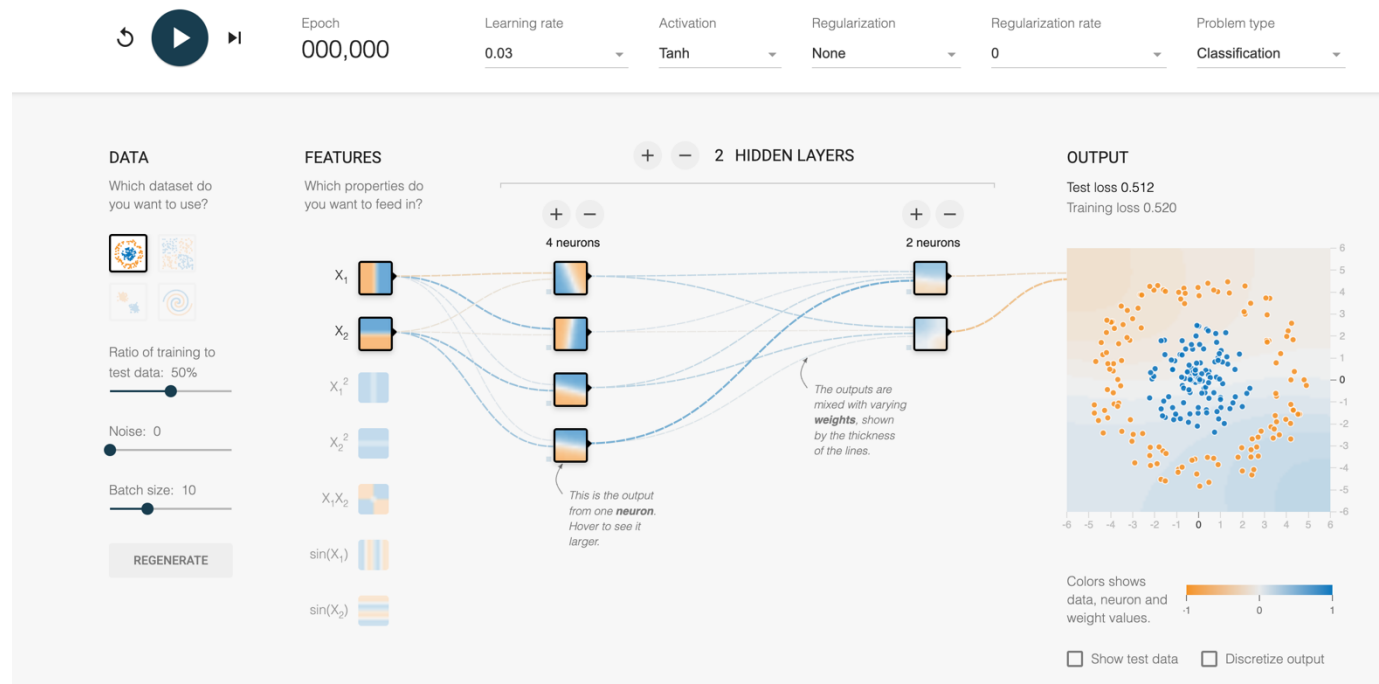
Layer number

Layer dimension

# A Neural Network Playground

Link to the [Neural Network Playground](#)

Tinker With a **Neural Network** Right Here in Your Browser.  
Don't Worry, You Can't Break It. We Promise.





# Training a Neural Network

# Loss Function

- Loss Function:  $L(f(x^{(i)}; W), y^{(i)})$ , where
  - $x^{(i)}$  are the inputs,  $f(\cdot; W)$  is the neural network
  - $f(x^{(i)}; W)$  is the predicted value
  - $y^{(i)}$  is the actual value
  - The loss of a neural network measures the cost incurred from incorrect predictions
- Empirical Loss:  $J(W) = \frac{1}{n} \sum_{i=1}^n L(f(x^{(i)}; W), y^{(i)})$ 
  - The empirical loss measures the total loss over the entire dataset
- Objective of training a neural network

$$W^* = \underset{W}{\operatorname{argmin}} J(W)$$



# Common Loss Functions

- Regression ( $f(x^{(i)}; W) \in \mathbb{R}$ ): **Mean squared error (MSE)**

$$J(W) = \frac{1}{n} \sum_{i=1}^n (y^{(i)} - f(x^{(i)}; W))^2$$

- Classification: **Cross-entropy loss**

- Binary classification ( $p^{(i)} = f(x^{(i)}; W) \in (0,1)$ ): Binary cross-entropy loss

$$J(W) = -\frac{1}{n} \sum_{i=1}^n [y^{(i)} \log(p^{(i)}) + (1 - y^{(i)}) \log(1 - p^{(i)})]$$

- Multi-class classification ( $p_c^{(i)} = f(x^{(i)}; W)_c \in (0,1)$ ): Categorical cross-entropy loss


$$J(W) = -\frac{1}{n} \sum_{i=1}^n \sum_{c=1}^C y_c^{(i)} \log p_c^{(i)}$$

- Multi-label classification: Each input belong to multiple classes simultaneously (e.g., image classification in self-driving)

$$J(W) = -\frac{1}{n} \sum_{i=1}^n \sum_{c=1}^C [y_c^{(i)} \log p_c^{(i)} + (1 - y_c^{(i)}) \log (1 - p_c^{(i)})]$$


# Optimization

- Objective of training a neural network
$$W^* = \underset{W}{\operatorname{argmin}} J(W)$$
- The loss function is a function of the network weights
- If we had only 2 weights, we could plot the loss function with respect to the different combinations of weights



# Gradient

- The **gradient** of  $J$  at the point  $W_0$ , denoted as  $\nabla_W J(W_0)$ , is the direction to move in for the fastest increase in  $J(W)$ , when starting from  $W_0$
- Thus, the opposite direction of the gradient points to the fastest decrease in  $J(W)$



# Gradient Descent

- Initialize  $W_0$  randomly
- For  $i$  from 0 to  $M$ :


$$W_{i+1} = W_i - \eta \nabla_W J(W_i)$$

  - $\eta$  is the step size, also called the learning rate
  - $M$  is the maximum number of iterations
- The algorithm continues until the stopping condition
  - When  $\| \nabla_W J(W) \|_2 \leq \varepsilon$  for some pre-set  $\varepsilon$  (Recall  $\nabla_W J(W) = 0$  when function  $J(W)$  is at minimum)
  - Evaluate the performance on validation data, and stop when not improving




# Learning Rate

- Step size or learning rate ( $\eta$ ) is a hyperparameter
- Small  $\eta$  leads to slow convergence, but large  $\eta$  leads to divergence
- Different step sizes have to be tried



# Backpropagation

Forward and backward pass: Chain rule



- Let's consider a simple case with two weight parameters,  $W = (w_1, w_2)$
- Gradient for  $w_2$

$$\frac{\partial J(w_1, w_2)}{\partial w_2} = \frac{\partial J(w_1, w_2)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w_2}$$

- Gradient for  $w_1$

$$\frac{\partial J(w_1, w_2)}{\partial w_1} = \frac{\partial J(w_1, w_2)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1} * \frac{\partial z_1}{\partial w_1}$$

- PyTorch has the autograd function that computes the gradient automatically
- YouTube video on autograd from scratch by Andrej Karpathy [[Link](#), Credit to Hitanshu]

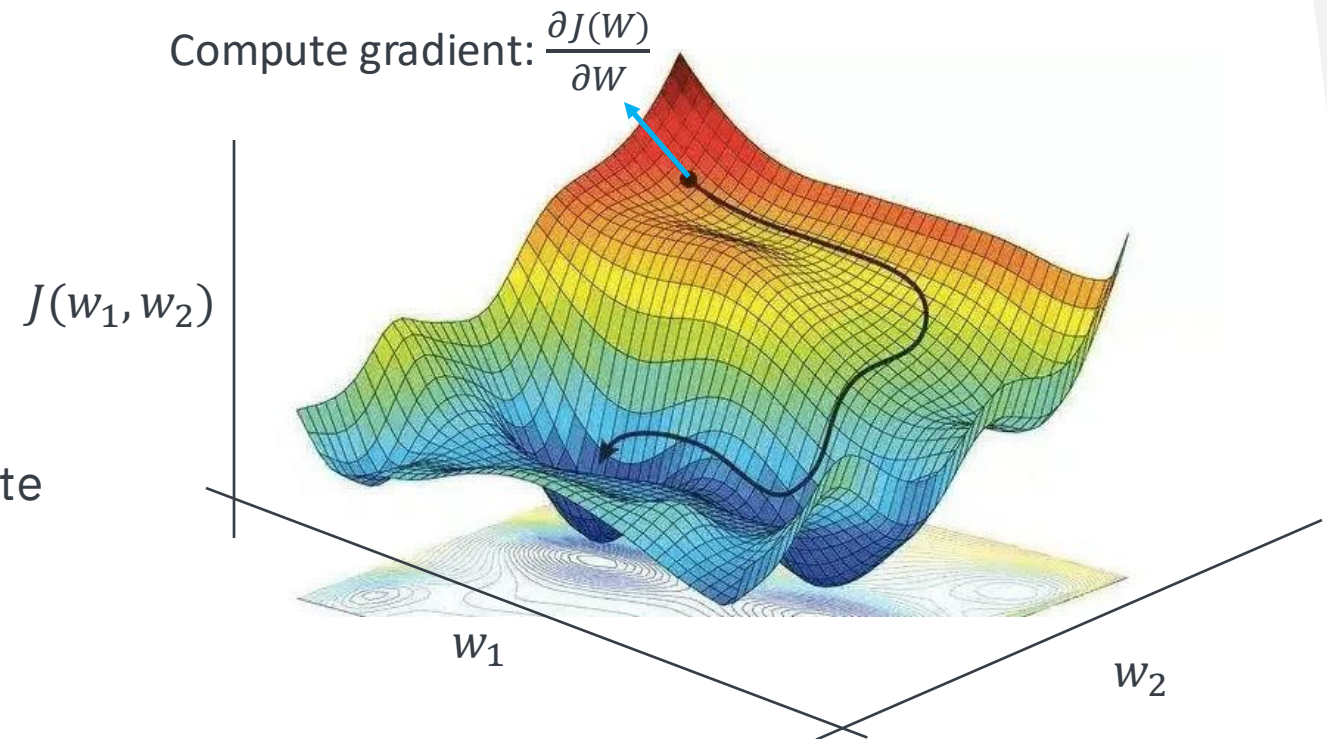


# Training Neural Networks with Gradient Descent

Objective:  $W^* = \underset{W}{\operatorname{argmin}} J(W)$ , where  $J(W) = \frac{1}{n} \sum_{i=1}^n L(f(x^{(i)}; W), y^{(i)})$

## Algorithm

1. Initialize weight matrix  $W$
2. Backpropagation: Loop until convergence
  - a) Forward pass: Compute the predictions and the loss
  - b) Backward pass: Compute the gradient with the chain rule
  - c) Update the weights with the learning rate
3. Return the weight matrix  $W$  and the predictions




The graph shows Energy ( $E$ ) on the vertical axis and weight ( $w$ ) on the horizontal axis. The function is non-convex, with multiple local minima. The global minimum is marked with a yellow dot. A local minimum is marked with a green dot. The region to the left of the local minimum is labeled 'NEGATIVE SLOPE' and the region to the right is labeled 'POSITIVE SLOPE'. A 'FLUCTUATION close to the LOCAL MINIMUM' is indicated by a green dot and arrows. The process of finding the global minimum is shown by decreasing  $w$  (moving left) and increasing  $w$  (moving right) from a local minimum.

- 
- STEVENS INSTITUTE *of* TECHNOLOGY


# Optimization with Batch

- Batch: A batch is a small subset of the training data used to update the model weights during training
  - Typically, we sample the batches at **random** (without replacement)
- Optimization with batch
  - Initialize  $W_0$  randomly
  - Compute gradient  $\nabla_W J^1(W_0)$   
Update  $W_1 = W_0 - \eta \nabla_W J^1(W_0)$
  - Compute gradient  $\nabla_W J^2(W_1)$   
Update  $W_2 = W_1 - \eta \nabla_W J^2(W_1)$
  - Compute gradient  $\nabla_W J^3(W_2)$   
Update  $W_3 = W_2 - \eta \nabla_W J^3(W_2)$
  - ...
- Epoch: An epoch is one complete pass through the entire training dataset
- Shuffle the batches after each epoch




# Small Batch vs. Large Batch

## MNIST



## CIFAR-10




- **MNIST** refers to the hand-written digit classification; **CIFAR-10** refers to the 10-class image classification
- Smaller batch size has better performance
- What's wrong with large batch size?




# Small Batch vs. Large Batch

Full Batch




Small Batch



# Momentum: Adaptive Learning Rate

Physical world: Movement has momentum



# Momentum: Adaptive Learning Rate

## Gradient Descent

Starting at  $\theta^0$


Compute gradient  $g^0$

Move to  $\theta^1 = \theta^0 - \eta g^0$

Compute gradient  $g^1$

Move to  $\theta^2 = \theta^1 - \eta g^1$

⋮



# Momentum: Adaptive Learning Rate

## Gradient Descent + Momentum

Starting at  $\theta^0$

Movement  $m^0 = 0$

Compute gradient  $g^0$

Movement  $m^1 = \lambda m^0 - \eta g^0$


Move to  $\theta^1 = \theta^0 + m^1$

Compute gradient  $g^1$

Movement  $m^2 = \lambda m^1 - \eta g^1$

Move to  $\theta^2 = \theta^1 + m^2$

Movement: **movement of last step** minus **gradient at present**





# Momentum: Adaptive Learning Rate

## Gradient Descent + Momentum

Starting at  $\theta^0$

Movement  $m^0 = 0$

Compute gradient  $g^0$

Movement  $m^1 = \lambda m^0 - \eta g^0$

Move to  $\theta^1 = \theta^0 + m^1$

Compute gradient  $g^1$

Movement  $m^2 = \lambda m^1 - \eta g^1$

Move to  $\theta^2 = \theta^1 + m^2$

Movement: **movement of last step** minus **gradient** at present

$m^i$  is the weighted sum of all the previous gradient:  $g^0, g^1, \dots, g^{i-1}$

$$m^0 = 0$$


$$m^1 = -\eta g^0$$

$$m^2 = -\lambda \eta g^0 - \eta g^1$$

$\vdots$

# Momentum: Adaptive Learning Rate

Gradient Descent + Momentum



# Local and Global Minimum

- Critical points have zero gradients
- Critical points can be either saddle points or local minima
- Smaller batch size and momentum help escape critical points
- **Adam** is a very popular optimizer that does all this automatically

## Adam: A method for stochastic optimization

[DP Kingma, J Ba](#) - arXiv preprint arXiv:1412.6980, 2014 - [arxiv.org](#)

... **Adam** works well in practice and compares favorably to other stochastic optimization methods.

Finally, we discuss AdaMax, a variant of **Adam** ... Overall, we show that **Adam** is a versatile ...




☆ Save 📄 Cite Cited by 202011 Related articles All 19 versions 🔗

```
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
optimizer = optim.Adam([var1, var2], lr=0.0001)
```



# Learning Curve


- A learning curve is a plot that shows the number of epochs on the x-axis and the loss values on the y-axis
- Learning curves of model performance on the train and validation datasets can be used to diagnose an underfit (left), overfit (right), or well-fit (middle) model



# Regularization


Dropout: Randomly dropping some nodes in the network during training to prevent the network from relying too much on a few nodes

- During training
  - In each forward pass during training, dropout randomly disables (or drops) a specified fraction of neurons in a layer by setting the neurons to zero
  - The remaining active neurons are scaled by a factor (often  $\frac{1}{1-\text{dropout rate}}$ ) to ensure that the output has the correct expected value
- During inference: When making predictions, dropout is not applied
- Dropout helps the network generalize by preventing it from relying too heavily on specific neurons, reducing the likelihood of overfitting



# Regularization

Early stopping: Stop training when performance on a validation set stops improving to prevent overfitting



# Regularization

L2-Regularization: Add a penalty to the loss function to reduce model complexity and prevent overfitting

- Similar to ridge regression

$$\min J(W) \longrightarrow \min J(W) + \lambda \cdot \sum_{i \in \text{layers}} \sum_{k \in \text{nodes of } i} W_{ki}^2$$

# Summary

Hyperparameters of a neural network that we need to tune

- The architecture of the neural network
  - The number of hidden layers
  - The number of hidden neurons in each hidden layer
  - The choice of the activation function
- Training a neural network
  - Batch size
  - Learning rate
  - Optimization
- Link to the [Neural Network Playground](#)





# Acknowledgement

The lecture note has benefited from various resources, including those listed below. Please contact Zonghao Yang (zyang99@stevens.edu) with any questions or concerns about the use of these materials.

- Lecture Notes on Deep Learning Fundamentals by Léonard Boussioux at University of Washington
- Lecture Notes on Deep Learning from ML 2021 Spring by Hung-Yi Lee at National Taiwan University





# THANK YOU

**Stevens Institute of Technology**  
1 Castle Point Terrace, Hoboken, NJ 07030